

# Bank\_Loans\_Modelling

September 3, 2020

## (1)Importing Library and Load data

```
In [3]: #importing required library
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import pandas.util.testing as tm
```

```
In [15]: #upload dataset from my local machine
from google.colab import files
uploaded = files.upload()
```

<IPython.core.display.HTML object>

Saving Bank\_Personal\_Loan\_Modelling.xlsx to Bank\_Personal\_Loan\_Modelling (1).xlsx

```
In [4]: #read excel file and store into a dataframe
df = pd.read_excel( 'Bank_Personal_Loan_Modelling.xlsx' , 'Data' )
```

```
In [5]: #used for calculating some statistical data like percentile, mean and std of the numer
df.describe()
```

```
Out[5]:
```

	ID	Age	...	Online	CreditCard
count	5000.000000	5000.000000	...	5000.000000	5000.000000
mean	2500.500000	45.338400	...	0.596800	0.294000
std	1443.520003	11.463166	...	0.490589	0.455637
min	1.000000	23.000000	...	0.000000	0.000000
25%	1250.750000	35.000000	...	0.000000	0.000000
50%	2500.500000	45.000000	...	1.000000	0.000000
75%	3750.250000	55.000000	...	1.000000	1.000000
max	5000.000000	67.000000	...	1.000000	1.000000

[8 rows x 14 columns]

```
In [6]: #show the starting five rows of the datasets
df.head()
```

```
Out [6]:
```

	ID	Age	Experience	...	CD Account	Online	CreditCard
0	1	25	1	...	0	0	0
1	2	45	19	...	0	0	0
2	3	39	15	...	0	0	0
3	4	35	9	...	0	0	0
4	5	35	8	...	0	0	1

[5 rows x 14 columns]

```
In [7]: #show the last five rows of the datasets
df.tail()
```

```
Out [7]:
```

	ID	Age	Experience	...	CD Account	Online	CreditCard
4995	4996	29	3	...	0	1	0
4996	4997	30	4	...	0	1	0
4997	4998	63	39	...	0	0	0
4998	4999	65	40	...	0	1	0
4999	5000	28	4	...	0	1	1

[5 rows x 14 columns]

```
In [8]: #show the non-null, count and datatypes
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    5000 non-null  int64
1   Age                   5000 non-null  int64
2   Experience             5000 non-null  int64
3   Income                5000 non-null  int64
4   ZIP Code              5000 non-null  int64
5   Family                5000 non-null  int64
6   CCAvg                 5000 non-null  float64
7   Education             5000 non-null  int64
8   Mortgage              5000 non-null  int64
9   Personal Loan         5000 non-null  int64
10  Securities Account    5000 non-null  int64
11  CD Account            5000 non-null  int64
12  Online                5000 non-null  int64
13  CreditCard            5000 non-null  int64
dtypes: float64(1), int64(13)
memory usage: 547.0 KB
```

```
In [9]: #show the shape of dataset
df.shape
```

```
Out[9]: (5000, 14)
```

```
In [10]: #this used to get an int representing the number of elements in this object  
df.size
```

```
Out[10]: 70000
```

```
In [19]: #check if there is null value or not and we found that Experience column contain null  
df.isnull().sum().sort_values(ascending = False)  
df = df.dropna(subset=['Experience'])  
  
df.isnull().sum().sort_values(ascending = False)
```

```
Out[19]: CreditCard      0  
         Online         0  
         CD Account     0  
         Securities Account 0  
         Personal Loan  0  
         Mortgage      0  
         Education     0  
         CCAvg         0  
         Family        0  
         ZIP Code      0  
         Income        0  
         Experience     0  
         Age           0  
         ID            0  
         dtype: int64
```

```
In [21]: #tells us about datatypes  
df.dtypes
```

```
Out[21]: ID            int64  
         Age           int64  
         Experience     int64  
         Income        int64  
         ZIP Code      int64  
         Family        int64  
         CCAvg         float64  
         Education     int64  
         Mortgage      int64  
         Personal Loan  int64  
         Securities Account int64  
         CD Account    int64  
         Online        int64  
         CreditCard    int64  
         dtype: object
```

```
In [22]: df.apply(lambda x : sum(x.isnull()))
```

```
Out[22]: ID          0
        Age          0
        Experience    0
        Income        0
        ZIP Code      0
        Family        0
        CCAvg         0
        Education     0
        Mortgage      0
        Personal Loan  0
        Securities Account 0
        CD Account    0
        Online        0
        CreditCard    0
        dtype: int64
```

## (2) Check if you need to clean the data for any of the variables

```
In [24]: #check that how many negative value is their for Expeience column , because Expeience
         df[df['Experience'] < 0]['Experience'].value_counts()
```

```
Out[24]: -1      33
        -2      15
        -3       4
        Name: Experience, dtype: int64
```

```
In [25]: #clean the negative variable
         dfExp = df.loc[df['Experience'] > 0]
         negExp = df.Experience < 0
         newlist = df.loc[negExp]['ID'].tolist() # getting the customer ID who has negative ex
         negExp.value_counts()
```

```
Out[25]: False      4948
        True         52
        Name: Experience, dtype: int64
```

```
In [26]: import numpy as np
```

```
In [27]: #Get the value of Age and Education columns
         #Filter the records which has positive experience and take the median
         #Apply the median to the location which had negative experience
```

```
In [28]: for id in newlist:
         age = df.loc[np.where(df['ID']==id)]["Age"].tolist()[0]
         df_filtered = dfExp[(dfExp.Age == age)]
         exp = df_filtered['Experience'].median()
         df.loc[df.loc[np.where(df['ID']==id)].index, 'Experience'] = exp
```

```
In [29]: df[df['Experience'] < 0]['Experience'].count()
```

```
Out[29]: 0
```

```
In [30]: df.describe()
```

```
Out[30]:
```

	ID	Age	...	Online	CreditCard
count	5000.000000	5000.000000	...	5000.000000	5000.000000
mean	2500.500000	45.338400	...	0.596800	0.294000
std	1443.520003	11.463166	...	0.490589	0.455637
min	1.000000	23.000000	...	0.000000	0.000000
25%	1250.750000	35.000000	...	0.000000	0.000000
50%	2500.500000	45.000000	...	1.000000	0.000000
75%	3750.250000	55.000000	...	1.000000	1.000000
max	5000.000000	67.000000	...	1.000000	1.000000

```
[8 rows x 14 columns]
```

### (3) Study the data distribution in each attribute and target variable

```
In [31]: import seaborn as sns
```

```
df.nunique() #Number of unique in each column
```

```
Out[31]: ID                5000
Age                  45
Experience            44
Income               162
ZIP Code             467
Family                4
CCAvg               108
Education             3
Mortgage             347
Personal Loan         2
Securities Account    2
CD Account            2
Online                2
CreditCard           2
dtype: int64
```

```
In [32]: zm = df[df['Mortgage'] == 0]['Mortgage'].count() #Number of people with zero mortgage
```

```
In [33]: #Number of people with zero credit card spending per month
df[df['CCAvg'] == 0]['CCAvg'].count()
```

```
Out[33]: 106
```

```
In [34]: df.count() #Value counts of all categorical columns
```

```
Out[34]: ID                5000
Age                  5000
```

Experience	4971
Income	5000
ZIP Code	5000
Family	5000
CCAvg	5000
Education	5000
Mortgage	5000
Personal Loan	5000
Securities Account	5000
CD Account	5000
Online	5000
CreditCard	5000

dtype: int64

```
In [35]: df.Family.value_counts()
```

```
Out[35]: 1    1472
         2    1296
         4    1222
         3    1010
         Name: Family, dtype: int64
```

```
In [36]: df.Education.value_counts()
```

```
Out[36]: 1    2096
         3    1501
         2    1403
         Name: Education, dtype: int64
```

```
In [37]: l=df.Online.value_counts()
```



```
In [38]: df.CreditCard.value_counts()
```


```
Out[38]: 0    3530
         1    1470
         Name: CreditCard, dtype: int64
```

```
In [39]: #Univariate analysis
```

```
In [40]: import matplotlib.pyplot as plt
import seaborn as sns
```



```
In [154]: plt.hist(df.Income, edgecolor='Red')
plt.xlabel('Income')
plt.ylabel('Customer')
plt.show()
plt.scatter(df.index,df['Income'])
plt.xlabel('Customer')
plt.ylabel('Income')
plt.show()
sns.boxplot(df['Income'])
plt.show()
```








```
In [42]: sns.distplot(df.Income)
plt.show()
sns.distplot(df.Age)
plt.show()
sns.distplot(df.CCAvg)
plt.show()
```










```
In [47]: plt.hist(df.Family, edgecolor='Red')
plt.xlabel('Family')
plt.ylabel('Customer')
plt.show()
plt.scatter(df.index,df['Family'])
plt.xlabel('Customer')
plt.ylabel('Family')
plt.show()
sns.boxplot(df['Family'])
plt.show()
```







```
In [48]: plt.hist(df['CCAvg'] , edgecolor='Red')
plt.xlabel('Avg. spending')
plt.show()
plt.scatter(df.index,df['CCAvg'])
plt.ylabel('Avg. spend')
plt.show()
```






```
In [49]: plt.hist(df.Education, edgecolor='Red')  
plt.xlabel('Education')  
plt.ylabel('Customer')
```

```
plt.show()
plt.scatter(df.index,df['Education'])
plt.xlabel('Customer')
plt.ylabel('Education')
plt.show()
```



```
In [50]: plt.hist(df.Mortgage, edgecolor='Red')
plt.xlabel('Mortgage')
plt.ylabel('Customer')
plt.show()
plt.scatter(df.index,df['Mortgage'])
plt.xlabel('Customer')
plt.ylabel('Mortgagae')
plt.show()
sns.boxplot(df['Mortgage'])
plt.show()
```







```
In [51]: plt.hist(df['Personal Loan'], edgecolor='Red')
plt.xlabel('Personal Loan')
```

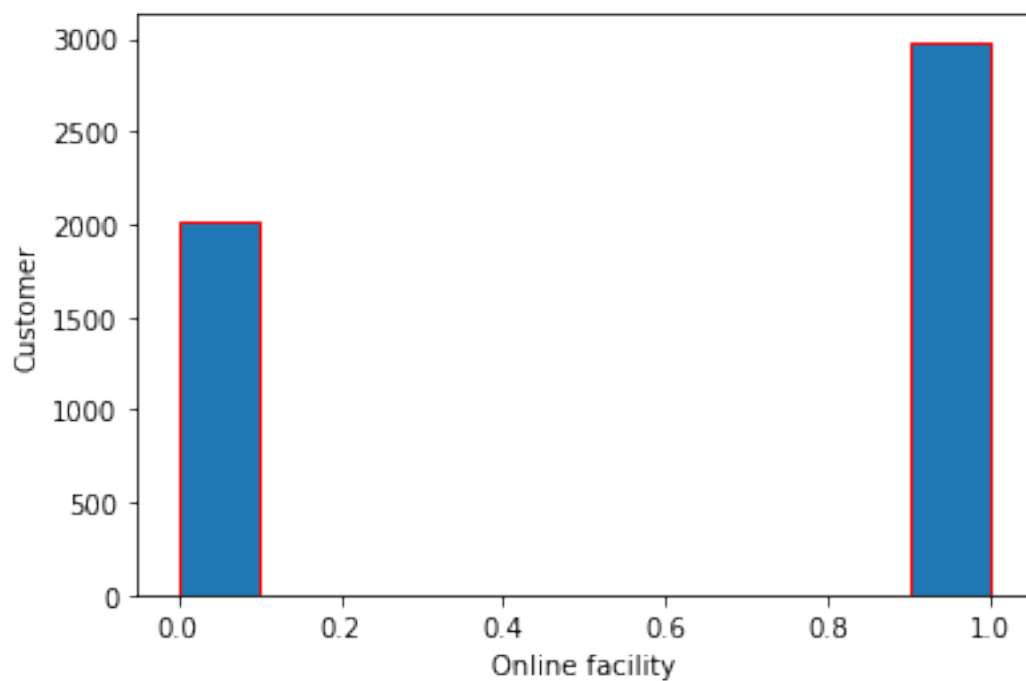


```
plt.ylabel('Customer')
plt.show()
plt.scatter(df.index,df['Personal Loan'])
plt.xlabel('Customer')
plt.ylabel('Personal Loan')
plt.show()
```






```
In [52]: plt.hist(df['Online'], edgecolor='Red')  
plt.xlabel('Online facility')  
plt.ylabel('Customer')  
plt.show()
```





```
In [53]: plt.hist(df['CreditCard'], edgecolor='Red')
plt.xlabel('Credit Card')
plt.ylabel('Customer')
plt.show()
```



```
In [54]: # Bivariate Analysis
```


```
In [55]: plt.scatter(df['CCAvg'], df['Income'])
plt.xlabel('Avg. spending')
plt.ylabel('Income')
df.plot.hexbin(x='CCAvg', y='CCAvg', gridsize=15)
```

```
Out[55]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9d6403fd30>
```




```
In [56]: plt.scatter(df['CCAvg'],df['Mortgage'])
plt.xlabel('Avg. spending')
plt.ylabel('Mortgage')
```

```
Out[56]: Text(0, 0.5, 'Mortgage')
```




```
In [57]: plt.scatter(df['Income'],df['Mortgage'])  
plt.xlabel('Income')  
plt.ylabel('Mortgage')
```

```
Out[57]: Text(0, 0.5, 'Mortgage')
```




```
In [58]: sns.boxplot(x='Education',y='Income',hue='Personal Loan',data=df)
```

```
Out[58]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9d645b1b00>
```



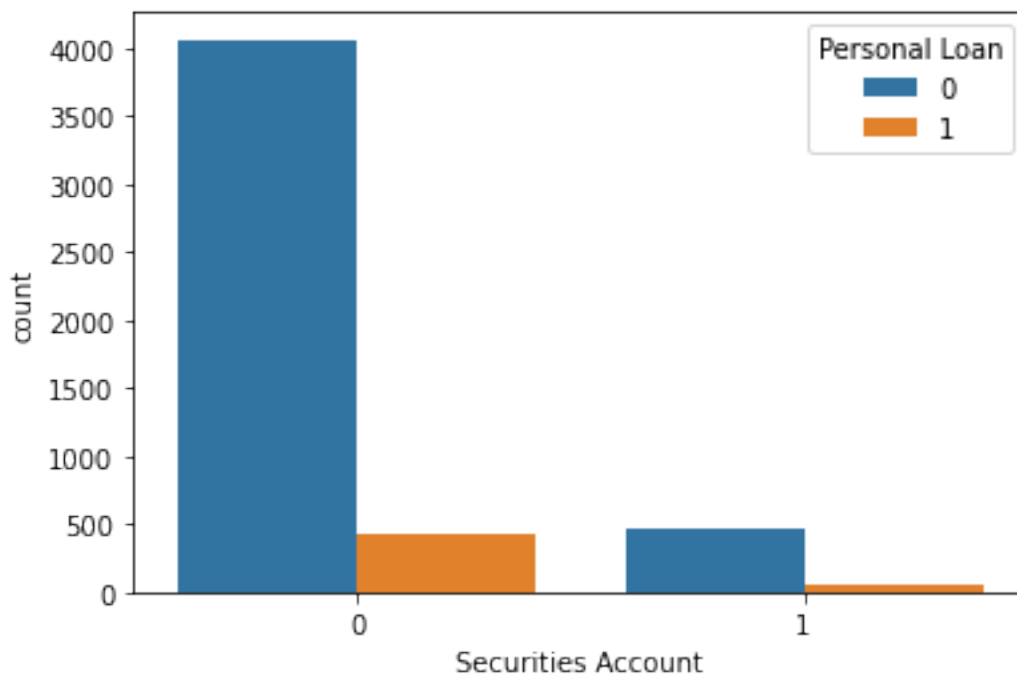
```
In [59]: sns.boxplot(x="Education", y='Mortgage', hue="Personal Loan", data=df)
```

```
Out[59]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9d64898dd8>
```




```
In [60]: sns.countplot(x="Securities Account", data=df, hue="Personal Loan")
```

```
Out[60]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9d63ed2f60>
```



```
In [61]: sns.countplot(x='CD Account',data=df,hue='Personal Loan')
```


```
Out[61]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9d647ba080>
```



```
In [62]: sns.distplot( df[df['Personal Loan'] == 0]['CCAvg'], color = 'r')
sns.distplot( df[df['Personal Loan'] == 1]['CCAvg'], color = 'g')
```

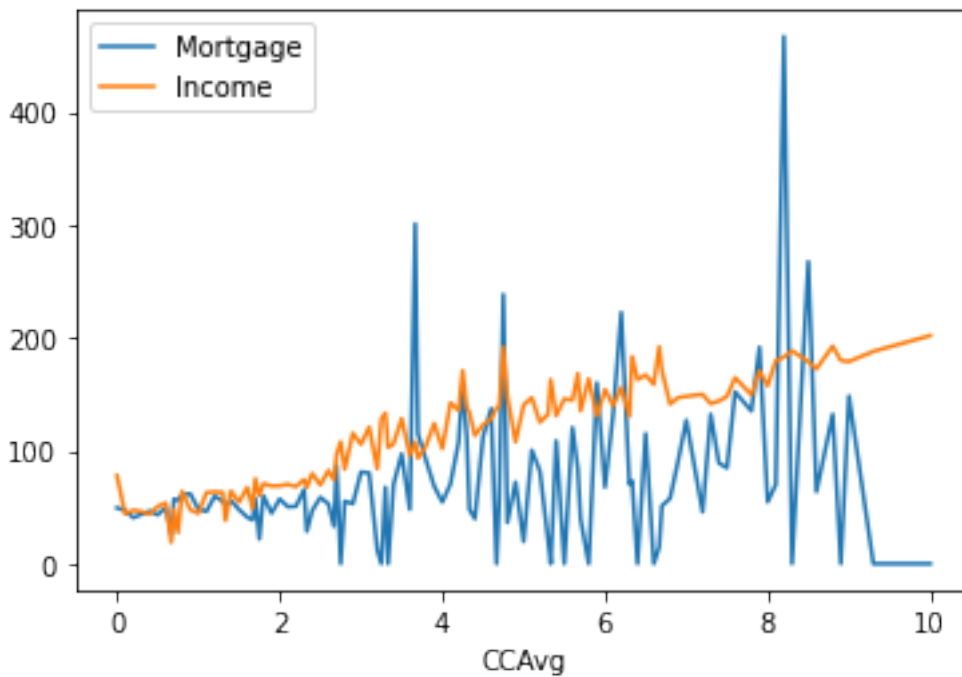
```
Out[62]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9d63d14b70>
```

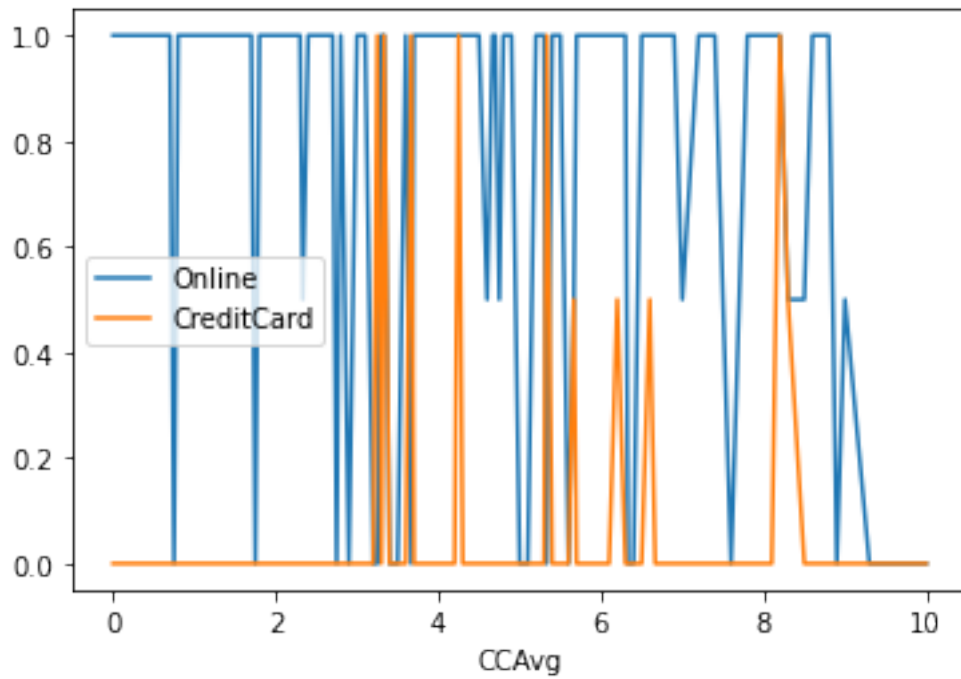




```
In [63]: df_group = df.groupby('CCAvg').mean()[['Mortgage', 'Income']]
df_group.plot.line()
df_group2 = df.groupby('CCAvg').median()[['Online', 'CreditCard']]
df_group2.plot.line()
```



Out[63]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f9d646d5be0>







```
In [64]: df_group3 = df.groupby('Personal Loan').mean()[['Mortgage', 'Income']]
df_group3.plot.line()
df_group3.plot.bar(stacked=True)
```

```
Out[64]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9d63e50cc0>
```




```
In [65]: df_group4 = df.groupby('CreditCard').mean()[['Online', 'Personal Loan', 'Securities Account']]
df_group4.plot.line()
df_group4.plot.bar(stacked=True)
```

Out[65]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f9d63c1e128>



```
In [66]: fig,ax = plt.subplots(figsize=(15,10))
sns.heatmap(df.corr() , cmap='plasma' , annot=True)
```

```
Out[66]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9d63b49518>
```



(5) Normalise your data and split the data into training and test set in the ratio of 70:30 respectively

```
In [67]: from sklearn.model_selection import train_test_split
target = 'Personal Loan'
df_x = df.drop(target,axis='columns',inplace=False)
df_y = df[target]

x_train,x_test,y_train,y_test = train_test_split(df_x, df_y,test_size=0.30,random_state=42)
```

(4) Apply necessary transformations for the feature variables

```
In [68]: def quick_analysis(df):
    print('Data Types:')
    print(df.dtypes)
    print('Rows and Columns:')
    print(df.shape)
```

```

print('Column Names:')
print(df.columns)
print('Null Values:')
print(df.apply(lambda x: sum(x.isnull()) / len(df)))

```

```

quick_analysis(df)

```

Data Types:

```

ID                int64
Age               int64
Experience         float64
Income            int64
ZIP Code          int64
Family            int64
CCAvg             float64
Education         int64
Mortgage          int64
Personal Loan     int64
Securities Account int64
CD Account        int64
Online            int64
CreditCard        int64

```

dtype: object

Rows and Columns:

(5000, 14)

Column Names:

```

Index(['ID', 'Age', 'Experience', 'Income', 'ZIP Code', 'Family', 'CCAvg',
      'Education', 'Mortgage', 'Personal Loan', 'Securities Account',
      'CD Account', 'Online', 'CreditCard'],
      dtype='object')

```

Null Values:

```

ID                0.0000
Age               0.0000
Experience         0.0058
Income            0.0000
ZIP Code          0.0000
Family            0.0000
CCAvg             0.0000
Education         0.0000
Mortgage          0.0000
Personal Loan     0.0000
Securities Account 0.0000
CD Account        0.0000
Online            0.0000
CreditCard        0.0000

```

dtype: float64

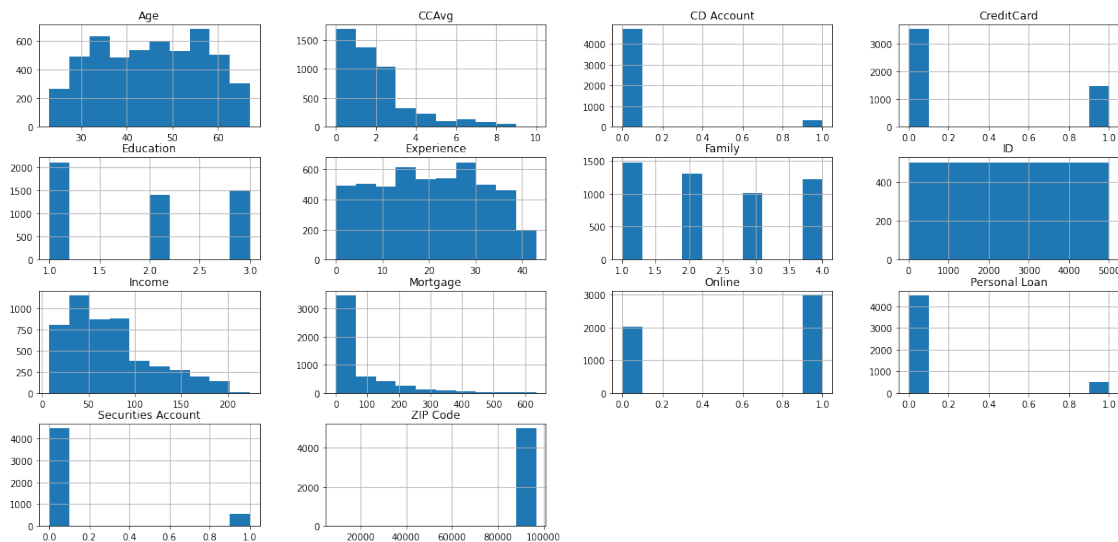
```
In [69]: df.describe()
```

```
Out[69]:
```

	ID	Age	...	Online	CreditCard
count	5000.000000	5000.000000	...	5000.000000	5000.000000
mean	2500.500000	45.338400	...	0.596800	0.294000
std	1443.520003	11.463166	...	0.490589	0.455637
min	1.000000	23.000000	...	0.000000	0.000000
25%	1250.750000	35.000000	...	0.000000	0.000000
50%	2500.500000	45.000000	...	1.000000	0.000000
75%	3750.250000	55.000000	...	1.000000	1.000000
max	5000.000000	67.000000	...	1.000000	1.000000


```
[8 rows x 14 columns]
```

```
In [70]: df[df.dtypes[(df.dtypes=="float64")|(df.dtypes=="int64")].index.values].hist(figsize=
plt.show())
```




```
In [71]: sns.distplot(df.Income)
```

```
Out[71]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9d5feb2780>
```



```
In [72]: sns.distplot(df.CCAvg)
```


```
Out[72]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9d5ca23f60>
```





```
In [73]: sns.distplot(df.Mortgage)
```


```
Out[73]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9d5c93e198>
```



```
In [74]: from sklearn.preprocessing import PowerTransformer as pt

pwt = pt(method="yeo-johnson" , standardize=False)

pwt.fit(df_x["Income"].values.reshape(-1,1))
temp = pwt.transform(df_x["Income"].values.reshape(-1,1))
sns.distplot(temp)
plt.show()
```



```
In [75]: from sklearn.preprocessing import PowerTransformer as pt
```


```
pwt = pt(method="yeo-johnson" , standardize=False)
```

```
pwt.fit(df_x["CCAvg"].values.reshape(-1,1))
```

```
temp = pwt.transform(df_x["CCAvg"].values.reshape(-1,1))
```

```
sns.distplot(temp)
```

```
Out[75]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9d5c7b9b70>
```



(6) Use the Logistic Regression model to predict the likelihood of a customer buying personal loans

```
In [76]: from sklearn.linear_model import LogisticRegression
         from sklearn.metrics import accuracy_score
         from sklearn.model_selection import train_test_split
```

```
In [84]: #In it our target is Personal Loan column so for x_dataset just drop this and for y_dataset
         target = 'Personal Loan'
         df_x = df.drop(target,axis='columns',inplace=False)
         df_y = df[target]
```

```
x_train,x_test,y_train,y_test = train_test_split(df_x, df_y,test_size=0.30,random_state=42)
```

```
In [79]: #check if their is null value exist
         df.isnull().sum().sort_values(ascending = False)
```

```
Out[79]: Experience      29
         CreditCard      0
         Online          0
         CD Account      0
         Securities Account  0
         Personal Loan    0
         Mortgage        0
         Education        0
         CCAvg            0
         Family           0
```

```

ZIP Code      0
Income        0
Age           0
ID            0
dtype: int64

```

```

In [82]: #if null value is their then drop
df = df.dropna(subset=['Experience'])
df.isnull().sum().sort_values(ascending = False)

```

```

Out[82]: CreditCard      0
Online      0
CD Account  0
Securities Account  0
Personal Loan  0
Mortgage    0
Education   0
CCAvg       0
Family      0
ZIP Code    0
Income      0
Experience   0
Age         0
ID          0
dtype: int64

```

```

In [140]: target = 'Personal Loan'
df_x = df.drop(target,axis='columns',inplace=False)
df_y = df[target]

```

```

#so perform Logistic Regression with the split of 70:30 in training and testing data
x_train,x_test,y_train,y_test = train_test_split(df_x, df_y,test_size=0.30,random_state=42)
L_R = LogisticRegression(max_iter=1000)
L_R.fit(x_train,y_train)

```

```

Out[140]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, l1_ratio=None, max_iter=1000,
                             multi_class='auto', n_jobs=None, penalty='l2',
                             random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                             warm_start=False)

```

```

In [122]: #Accuracy for the model
y_pred = L_R.predict(x_test)
print("Accuracy of logistic regression classifier on test set :",accuracy_score(y_test,y_pred))

```

```

Accuracy of logistic regression classifier on test set : 94.16890080428955

```

```

In [123]: print("Testing Accuracy",L_R.fit(x_train,y_train).score(x_test,y_test)*100)

```

Testing Accuracy 94.16890080428955

```
In [124]: print("Training Accuracy" , L_R.fit(x_train,y_train).score(x_train,y_train)*100)
```

Training Accuracy 93.7625754527163

### (7) Print all the metrics related for evaluating the model performance

```
In [125]: #classification report contains all major aspect which are the very useful for analy
          from sklearn.metrics import classification_report
          print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.96	0.98	0.97	1348
1	0.76	0.58	0.66	144
accuracy			0.94	1492
macro avg	0.86	0.78	0.81	1492
weighted avg	0.94	0.94	0.94	1492

```
In [126]: from sklearn.metrics import confusion_matrix
          confusion_matrix = confusion_matrix(y_test, y_pred)
          print(confusion_matrix)
```

```
[[1321  27]
 [  60  84]]
```

```
In [138]: #show the confusion matrix with and without normalization
          from sklearn.metrics import plot_confusion_matrix
          classifier = LogisticRegression(max_iter=1000,C=0.01).fit(x_train, y_train)

          np.set_printoptions(precision=2)

          # Plot non-normalized confusion matrix
          titles_options = [("Confusion matrix, without normalization", None),
                           ("Normalized confusion matrix", 'true')]
          for title, normalize in titles_options:
              disp = plot_confusion_matrix(classifier, x_test, y_test,
                                           cmap=plt.cm.Red,
                                           normalize=normalize)

              disp.ax_.set_title(title)

              print(title)
```

```
print(dispatch.confusion_matrix)
```

```
plt.show()
```

Confusion matrix, without normalization


```
[[1326  22]
```


```
 [ 69  75]]
```

Normalized confusion matrix

```
[[0.98 0.02]
```


```
 [0.48 0.52]]
```





In [153]: *#ROC ( Receiver Operating Characteristic ) Curve*

```
#a good classifier stays as far away from that line as possible (toward the top-left
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
logit_roc_auc = roc_auc_score(y_test, L_R.predict(x_test))
fpr, tpr, thresholds = roc_curve(y_test, L_R.predict_proba(x_test)[: ,1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' % logit_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



```
In [129]: pred_1_test_x = L_R.predict(x_test)
          pred_1_train_x = L_R.predict(x_train)

          from sklearn.metrics import recall_score,precision_score,f1_score,roc_auc_score
          print("Recall : ",recall_score(y_test , pred_1_test_x))
          print("Precision : ", precision_score(y_test , pred_1_test_x))
          print("F1 Score : ", f1_score(y_test , pred_1_test_x))
          print("Roc Auc Score : ", roc_auc_score(y_test , pred_1_test_x))
```

*#Normal Logistic Regression give the accuracy 94% with 58% Recall and 78% Roc.*

```
Recall : 0.5833333333333334
Precision : 0.7567567567567568
F1 Score : 0.6588235294117648
Roc Auc Score : 0.7816518298714146
```

## (8) Build various other classification algorithms and compare their performance

```
In [130]: #AFTER APPLYING STANDARDIZATION IN LOGISTIC REGRESSION
```

```
from sklearn import preprocessing

col_names = df.columns
```



```

scaler=preprocessing.StandardScaler()
scaled_x_train=scaler.fit_transform(x_train)
scaled_x_test=scaler.fit_transform(x_test)
L_R = LogisticRegression()
L_R.fit(scaled_x_train,y_train)

from sklearn.metrics import recall_score,precision_score,f1_score,roc_auc_score,accuracy_score
from sklearn.metrics import roc_curve,auc

y_pred1 = L_R.predict(scaled_x_test)
print(classification_report(y_test,y_pred1))
print(accuracy_score(y_test,y_pred1))
print(confusion_matrix(y_test,y_pred1))

LR_prob=L_R.predict_proba(scaled_x_test)
fpr1,tpr1,thresholds1=roc_curve(y_test,LR_prob[:,1])
roc_auc1=auc(fpr1,tpr1)
print("Area under the Roc curve : %f " % roc_auc1)

pred_1_test_x = L_R.predict(scaled_x_test)
pred_1_train_x = L_R.predict(scaled_x_train)
print("Recall : ",recall_score(y_test , pred_1_test_x))
print("Precision : ", precision_score(y_test , pred_1_test_x))
print("F1 Score : ", f1_score(y_test , pred_1_test_x))
print("Roc Auc Score : ", roc_auc_score(y_test , pred_1_test_x))

# Standardization Logistic Regression give the accuracy 95% with 63% Recall and 81% Precision
# which is better then Logistic Regression .

```

	precision	recall	f1-score	support
0	0.96	0.98	0.97	1348
1	0.81	0.64	0.72	144
accuracy			0.95	1492
macro avg	0.89	0.81	0.84	1492
weighted avg	0.95	0.95	0.95	1492

```

0.9510723860589813
[[1327  21]
 [ 52  92]]
Area under the Roc curve : 0.950137
Recall : 0.6388888888888888
Precision : 0.8141592920353983
F1 Score : 0.7159533073929961
Roc Auc Score : 0.811655126937026

```

In [131]: *#K-Nearest Neighbor*

```

from sklearn.metrics import roc_curve,auc
from sklearn.neighbors import KNeighborsClassifier
knn_Model=KNeighborsClassifier(n_neighbors=3)
knn_Model.fit(scaled_x_train,y_train)

y_pred=knn_Model.predict(scaled_x_test)

print(classification_report(y_test,y_pred))
print(accuracy_score(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))

kNN_prob=knn_Model.predict_proba(scaled_x_test)
fpr2,tpr2,thresholds2=roc_curve(y_test,kNN_prob[:,1])
roc_auc2=auc(fpr2,tpr2)
print("Area under the Roc curve : %f " % roc_auc2)

pred_1_test_x = knn_Model.predict(scaled_x_test)
pred_1_train_x = knn_Model.predict(scaled_x_train)
print("Recall : ",recall_score(y_test , pred_1_test_x))
print("Precision : ", precision_score(y_test , pred_1_test_x))
print("F1 Score : ", f1_score(y_test , pred_1_test_x))
print("Roc Auc Score : ", roc_auc_score(y_test , pred_1_test_x))

# K-Nearest Neighbor give the accuracy 94% with 53% Recall and 76% Roc.
# which is worse then Normal Logistic Regression.

```

	precision	recall	f1-score	support
0	0.95	0.99	0.97	1348
1	0.88	0.53	0.66	144
accuracy			0.95	1492
macro avg	0.91	0.76	0.82	1492
weighted avg	0.94	0.95	0.94	1492

```

0.9477211796246648
[[1337  11]
 [ 67  77]]
Area under the Roc curve : 0.884525
Recall : 0.5347222222222222
Precision : 0.875
F1 Score : 0.6637931034482758
Roc Auc Score : 0.7632809924167491

```

```

In [132]: #NAIVE BAYES
from sklearn.naive_bayes import GaussianNB

```

```

naive_model = GaussianNB()
naive_model.fit(scaled_x_train,y_train)
y_pred=naive_model.predict(scaled_x_test)

print(classification_report(y_test,y_pred))
print(accuracy_score(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))

nbm_prob=naive_model.predict_proba(scaled_x_test)
fpr3,tpr3,thresholds3=roc_curve(y_test,nbm_prob[:,1])
roc_auc3=auc(fpr3,tpr3)
print("Area under the Roc curve : %f " % roc_auc3)

pred_1_test_x = naive_model.predict(scaled_x_test)
pred_1_train_x = naive_model.predict(scaled_x_train)
print("Recall : ",recall_score(y_test , pred_1_test_x))
print("Precision : ", precision_score(y_test , pred_1_test_x))
print("F1 Score : ", f1_score(y_test , pred_1_test_x))
print("Roc Auc Score : ", roc_auc_score(y_test , pred_1_test_x))

# Naive Bayes give the accuracy 88% with 53% Recall and 72% Roc.
# which is even worse then K-Nearest Neighbor.

```

	precision	recall	f1-score	support
0	0.95	0.92	0.93	1348
1	0.41	0.53	0.46	144
accuracy			0.88	1492
macro avg	0.68	0.73	0.70	1492
weighted avg	0.90	0.88	0.89	1492

```

0.8806970509383378
[[1237 111]
 [ 67  77]]
Area under the Roc curve : 0.920989
Recall : 0.5347222222222222
Precision : 0.4095744680851064
F1 Score : 0.463855421686747
Roc Auc Score : 0.7261890042861853

```

```

In [133]: #SUPPORT VECTOR MACHINE
from sklearn import svm
clf=svm.SVC(C=3, kernel='rbf', probability=True)

```

```

clfr.fit(scaled_x_train,y_train)
y_pred= clfr.predict(scaled_x_test)

print(classification_report(y_test,y_pred))
print(accuracy_score(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))

svm_prob=clfr.predict_proba(scaled_x_test)
fpr4,tpr4,thresholds4=roc_curve(y_test,svm_prob[:,1])
roc_auc4=auc(fpr4,tpr4)
print("Area under the Roc curve : %f " % roc_auc4)

pred_1_test_x = clfr.predict(scaled_x_test)
pred_1_train_x = clfr.predict(scaled_x_train)
print("Recall : ",recall_score(y_test , pred_1_test_x))
print("Precision : ", precision_score(y_test , pred_1_test_x))
print("F1 Score : ", f1_score(y_test , pred_1_test_x))
print("Roc Auc Score : ", roc_auc_score(y_test , pred_1_test_x))

# Support Vector Machine give the accuracy 97% with 80% Recall and 90% Roc.
# which is better then Standardization Logistic Regression.

```

	precision	recall	f1-score	support
0	0.98	1.00	0.99	1348
1	0.96	0.81	0.88	144
accuracy			0.98	1492
macro avg	0.97	0.90	0.93	1492
weighted avg	0.98	0.98	0.98	1492

```

0.9778820375335121
[[1343    5]
 [ 28 116]]
Area under the Roc curve : 0.981155
Recall : 0.8055555555555556
Precision : 0.9586776859504132
F1 Score : 0.8754716981132076
Roc Auc Score : 0.9009231783712496

```

In [134]: *#DECISION TREE*

```

from sklearn.tree import DecisionTreeClassifier
d_t=DecisionTreeClassifier(criterion='entropy' , random_state=1)
d_t.fit(scaled_x_train,y_train)
y_pred=d_t.predict(scaled_x_test)

```

```

print(classification_report(y_test,y_pred))
print(accuracy_score(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))

dt_prob=d_t.predict_proba(scaled_x_test)
fpr5,tpr5,thresholds5=roc_curve(y_test,dt_prob[:,1])
roc_auc5=auc(fpr5,tpr5)
print("Area under the Roc curve : %f " % roc_auc5)

pred_1_test_x = d_t.predict(scaled_x_test)
pred_1_train_x = d_t.predict(scaled_x_train)
print("Recall : ",recall_score(y_test , pred_1_test_x))
print("Precision : ", precision_score(y_test , pred_1_test_x))
print("F1 Score : ", f1_score(y_test , pred_1_test_x))
print("Roc Auc Score : ", roc_auc_score(y_test , pred_1_test_x))

# Decison Tree give the accuracy 98.3% with 87% Recall and 93.4% Roc.
# which is better then SVM.

```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	1348
1	0.95	0.88	0.91	144
accuracy			0.98	1492
macro avg	0.97	0.93	0.95	1492
weighted avg	0.98	0.98	0.98	1492

```

0.9832439678284183
[[1341    7]
 [ 18 126]]
Area under the Roc curve : 0.934904
Recall : 0.875
Precision : 0.9473684210526315
F1 Score : 0.9097472924187725
Roc Auc Score : 0.9349035608308606

```

In [135]: *#RANDOM FOREST*

```

from sklearn.ensemble import RandomForestClassifier
r_f=RandomForestClassifier(criterion='entropy',n_estimators=150,max_features=6,random
r_f.fit(scaled_x_train,y_train)
y_pred=r_f.predict(scaled_x_test)

print(classification_report(y_test,y_pred))
print(accuracy_score(y_test,y_pred))
print(confusion_matrix(y_test,y_pred))

```

```

rf_prob=r_f.predict_proba(scaled_x_test)
fpr6,tpr6,thresholds6=roc_curve(y_test,rf_prob[:,1])
roc_auc6=auc(fpr6,tpr6)
print("Area under the Roc curve : %f " % roc_auc6)

pred_1_test_x = r_f.predict(scaled_x_test)
pred_1_train_x = r_f.predict(scaled_x_train)
print("Recall : ",recall_score(y_test , pred_1_test_x))
print("Precision : ", precision_score(y_test , pred_1_test_x))
print("F1 Score : ", f1_score(y_test , pred_1_test_x))
print("Roc Auc Score : ", roc_auc_score(y_test , pred_1_test_x))

# Random Forest give the accuracy 98.7% with 86% Recall and 93% Roc.
# which is better then Decision Tree.

```

	precision	recall	f1-score	support
0	0.99	1.00	0.99	1348
1	1.00	0.87	0.93	144
accuracy			0.99	1492
macro avg	0.99	0.93	0.96	1492
weighted avg	0.99	0.99	0.99	1492

```

0.9872654155495979
[[1348    0]
 [  19  125]]
Area under the Roc curve : 0.997306
Recall :  0.8680555555555556
Precision :  1.0
F1 Score :  0.929368029739777
Roc Auc Score :  0.9340277777777778

```

## (9) Business understanding of the model

In [137]: *#It show the comparison boxplot graph of various machinelearning algorithms model*

```

from sklearn import model_selection

models = []
models.append(('KNN', KNeighborsClassifier()))
models.append(('LR', LogisticRegression(max_iter=1000)))
models.append(('DT', DecisionTreeClassifier()))
models.append(('NB', GaussianNB()))
models.append(('RF', RandomForestClassifier()))

```

```


# evaluate each model in turn
results = []
names = []
scoring = 'accuracy'
for name, model in models:
    kfold = model_selection.KFold(n_splits=10)
    cv_results = model_selection.cross_val_score(model, df_x, df_y, cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)
# boxplot algorithm comparison
fig = plt.figure()
fig.suptitle('Algorithm Comparison')
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()

```

```

KNN: 0.899819 (0.015683)
LR: 0.913297 (0.011714)
DT: 0.979285 (0.010395)
NB: 0.887546 (0.009978)
RF: 0.986522 (0.003123)

```



## Conclusion

The aim of the bank is to convert there liability customers into loan customers. They want to set up a new marketing campaign; hence, they need information about the connection between the variables given in the data. So from six classification algorithms were used in this study. From the above graph , it seems like **Random Forest algorithm** have the highest accuracy and we can choose that as our final model.